

The Longest Path Problem

Computational Complexity Analysis

Longest Path Problem

Input: Weighted graph $G = (V, E, w)$ where $w : E \rightarrow \mathbb{R}^+$

```
8 15 # Vertices, Edges  
v2 v1 5 # vertex, vertex, weight  
...
```

Output: A simple path P with maximum sum of edge weights

```
115  
v5 v4 v7 v8 v3 v1 v6 v2
```

Optimization Version (NP-Hard): Given a weighted graph, find the simple path with maximum total weight.

Decision Version (NP-Complete): Given a weighted graph, does a simple path of weight $\geq k$ exist?

Shortest vs Longest Path

Why can't we use a shortest path algorithm like Bellman-Ford (Polynomial) to solve Longest Path?

Attempt #1 - Negate every weight:

Definitionally, Shortest Path is *not* required to be simple, Longest Path *is*.

- Bellman-Ford will stop if it encounters a negative weight cycle to prevent an infinite loop.
- Even in the presence of a positive weight cycle, vertices in Longest Path can not be repeated and so it will never loop.

Note: Allowing repeated vertices and positive weight cycle detection in Longest Path would make it polynomial. Conversely, requiring Shortest Path to be simple would make it NP-Hard.

Shortest vs Longest Path

Why can't we use a shortest path algorithm like Bellman-Ford (Polynomial) to solve Longest Path?

Attempt #2 - Use reciprocal weights:

This doesn't work either because the reciprocal of a sum is not equal to the sum of reciprocals (e.g. $\frac{1}{2+3} \neq \frac{1}{2} + \frac{1}{3}$)

Counter-Example

- Path A: $w_1 = 1, w_2 = 4$
- Path B: $w_3 = 2, w_4 = 2$

Longest Path: $(1 + 4) > (2 + 2) \therefore$ Longest Path: (w_1, w_2)

Shortest Path + Reciprocals: $(1 + \frac{1}{4}) > (\frac{1}{2} + \frac{1}{2}) \therefore$ Longest Path: (w_3, w_4)

NP (Certificate Verification)

Certificate: A sequence of vertices $P = [v_1, v_2, \dots, v_m]$ claimed to be a longest path

```
def verify_longest_path(G, certificate, k):
    # Check 1: O(|P|) - All vertices in certificate are unique (simple path)
    if len(certificate) != len(set(certificate)):
        return False

    # Check 2: O(|P|) - All consecutive pairs form edges in G
    total_weight = 0 # Accumulator
    for i in range(len(certificate) - 1):
        edge = G.get_edge(certificate[i], certificate[i+1])
        if edge is None: # Verify edge exists in graph
            return False

        total_weight += edge.weight

    # Check 3: O(1) - Total path weight ≥ k
    return total_weight >= k
```


NP-Hardness

Hamiltonian Path Problem (known NP-Complete):

- Input: Unweighted Graph $G = (V, E)$
- Output: A path visiting every vertex exactly once (if it exists)

Reduction Strategy: Hamiltonian Path \leq_p Longest Path

If the Hamiltonian Path Problem can be reduced to Longest Path in polynomial time, then the Longest Path problem is also NP-Hard.



Reduction

Reduce Hamiltonian Path \rightarrow Longest Path in polynomial time:

Assign every edge in the graph a weight of 1.


$G = (V, E) \rightarrow G = (V, E, w)$ where $w : E \rightarrow 1$.

$O(|E|)$

Reframe decision:

Does a path of total weight $\geq |V| - 1$ exist?

- If YES, then the path found by Longest Path is a Hamiltonian path.
- If NO, then no Hamiltonian Path exists.



Reduction Correctness

Forward (\Rightarrow): If G has Hamiltonian path P :

- P visits all vertices exactly once
- P uses exactly $|V| - 1$ edges.
- All edges have weight 1 so the total weight of P is $|V| - 1$
- Thus there exists a simple path of total weight $\geq |V| - 1$, so Longest Path answers YES ✓

Reduction Correctness

Backward (\Leftarrow): If Longest Path returns path P of weight $\geq |V| - 1$:

- P is simple (no repeated vertices)
- P uses exactly $|V|$ vertices in a graph with $|V|$ total vertices
- Therefore P visits every vertex exactly once
- P is a Hamiltonian path ✓

Conclusion: Reduction runs in polynomial time, therefore, because Hamiltonian Path is NP-Complete, Longest Path is NP-Hard

Implementation (Simplified)

```
# graph: dict[int, dict[int, float]]\n\ndef recurse(\n    start: int, weight: float, visited: list[bool]\n) -> tuple[list[int], float]:\n    # ...\n\n    for adj in graph[start].keys():\n        if visited[adj]:\n            continue\n        # DOMINANT OPERATION: Recursive exploration\n        new_path, new_weight = recurse(adj, weight + graph[start][adj], visited)\n        # ...\n\n        if new_weight > best_weight:\n            best_weight = new_weight\n            new_path.insert(0, start)\n            best_path = new_path\n\n    return (best_path, best_weight)
```

Analytical Runtime Analysis

Worst Case (Complete Graph) Analysis:

- Start from each of n vertices: $O(n)$
- At each vertex, try extending to each unvisited neighbor
- At depth d , we have d visited vertices, $(n - d)$ unvisited
- Total simple paths explored: $n \cdot (n - 1) \cdot (n - 2) \cdot \dots \cdot 1 = n!$
- Total: $O(n! \cdot n) = O(n!) = O(|V|!)$

Dominant Operation: Recursive call exploring all possible paths


This factorial runtime is the dominant term and confirms the exponential nature of the problem

Empirical Runtime Analysis

Test Setup: Test cases timed with zsh time built-in

Test Case	Vertices	Edges	Time - Python	Time - Rust	Test Case	Vertices	Edges	Time - Python	Time - Rust
lp_172.txt	15	28	0:00.159	0:00.004	lp_136.txt	19	43	0:13.131	0:00.174
lp_083.txt	11	36	0:01.509	0:00.027	lp_175.txt	17	45	0:34.522	0:00.429
lp_156.txt	20	38	0:01.195	0:00.020	lp_135.txt	18	50	1:06.896	0:00.720
lp_100.txt	12	40	0:05.613	0:00.080	lp_151.txt	21	54	3:55.820	0:02.478
lp_152.txt	16	41	0:08.921	0:00.119	lp_108.txt	21	65	DNF (>20min)	2:02.490
lp_101.txt	16	42	0:12.165	0:00.177					

Exact Solution Runtime



High-Level Approximation Strategy

The algorithm combines **two major ideas**:

1. Greedy
2. Random Exploration

Greedy Construction

```
def GREEDY_STEP(G, current, visited):
    # All neighbors of current that are not yet visited
    candidates = [v for v in G[current] if v not in visited]
    if not candidates:
        return None

    # Score(v): degree + number of unvisited neighbors
    def SCORE(v):
        degree = len(G[v])
        unseen = sum(1 for u in G[v] if u not in visited)
        return (degree, unseen)

    # Return the candidate with maximum score
    return max(candidates, key=SCORE)
```

Random Exploration

```
def RANDOM_EXPLORE(G, current, visited, JUMP_PROB):
    r = random.random()      # number in [0, 1)

    # With probability JUMP_PROB, explore randomly
    if r < JUMP_PROB:
        candidates = [v for v in G[current] if v not in visited]
        if not candidates:
            return None
        return random.choice(candidates)

    # Otherwise, follow greedy rule
    return GREEDY_STEP(G, current, visited)
```

Runtime Analysis of `sample_path`

We analyze the worst-case runtime of one greedy + random-jump path.

At each step:

- Scan all neighbors of the current vertex → **O(degree(current))**
- Filter unvisited neighbors
- Pick best-weight or random-choice → **O(1)**

Each vertex becomes `current` at most once,
so each edge is scanned at most **twice**.

Overall Worst-Case Runtime


Total neighbor scanning:

$$\begin{aligned} & \text{degree}(v_1) + \text{degree}(v_2) + \dots + \text{degree}(v_n) \\ &= 2E = O(E) \end{aligned}$$


Thus:

$$T_{\text{sample}}(E) = O(E)$$


Each sampled path runs in
linear time in the number of edges.




Exact vs Approximation Longest-Path Values (Time = 0.001s)




Exact vs Approximation Longest-Path Values (Time = 0.01s)




Exact vs Approximation Longest-Path Values (Time = 0.1s)



Exact vs Approximation Longest-Path Values (Time = 1s)



Exact vs Approximation Longest-Path Values (Time = 5s)



Anytime Approximation Strategy

The algorithm has **two main components**:

1. Greedy Top-k Selection
2. Random Edge Selection

Greedy Top-k Selection (Pseudocode)

```
# Only consider unvisited neighbors
candidates = [(neighbor, weight)]
    for neighbor, weight in graph[current_vertex].items()
        if neighbor not in visited_vertices]

if not candidates:
    break

# pick randomly from top-k heaviest using heapq.nlargest
top_k_candidates = heapq.nlargest(k, candidates, key=lambda x: x[1])
neighbor, weight = random.choice(top_k_candidates)

visited_vertices.add(neighbor)
current_path.append(neighbor)
current_weight += weight
current_vertex = neighbor
```

Random Edge Selection (Pseudocode)

```
# Only consider unvisited neighbors
candidates = [(neighbor, weight)]
    for neighbor, weight in graph[current_vertex].items()
        if neighbor not in visited_vertices]

if not candidates:
    break

# probability set to 0.3 by default
if random.random() <= probability:
    neighbor, weight = random.choice(candidates)
else:
    # pick heaviest neighbor in linear time using max
    neighbor, weight = max(candidates, key=lambda x: x[1])

visited_vertices.add(neighbor)
current_path.append(neighbor)
current_weight += weight
current_vertex = neighbor
```

Runtime Analysis of the Algorithm

Let:

n = number of vertices

E = number of edges

Δ = max degree of a vertex

Top-K Selection Runtime

- **Per vertex iteration:**
 - Scan neighbors: $O(\Delta)$ (Δ = degree of current vertex)
 - Pick top-k using `heapq.nlargest`: $O(\Delta)$
 - Random choice from top-k: $O(1)$
- **Total Runtime:**
 - Worst-case $O(\Delta) = O(E)$

Random Edge Selection Runtime

- **Per vertex iteration:**
 - Scan neighbors: $O(\Delta)$
 - Random pick: $O(1)$ or max by weight: $O(\Delta)$
- **Total Runtime:**
 - Worst-case $O(\Delta) = O(E)$

Anytime Algorithm Total Runtime

- Top-k Selection $O(E)$ + Random Edge Selection $O(E) = O(E)$
- Algorithm repeats these iterations until time limit is reached,
- **Total Runtime Per Iteration:** $O(E)$
- **Total Runtime:** unbounded (anytime nature)
 - Can run as long as needed, producing progressively better solutions

Calculating an Upper Bound


Use a **Maximum Spanning Tree**:

- A spanning tree connects all n vertices using exactly $n-1$ edges
- An MST is the spanning tree with maximum total weight among all trees
- The **Longest Simple Path Weight \leq Maximum Spanning Tree Weight**

Patrick's Runtime VS Alston's Runtime

FILE	UPPER BOUND	1 SECOND	2.5 SECONDS	10 SECONDS	30 SECONDS	60 SECONDS	FILE	UPPER BOUND	1 SECOND	2.5 SECONDS	10 SECONDS	30 SECONDS	60 SECONDS
lp_151	315	260	264	265	265	265	lp_151	315	264	265	265	265	265
lp_184	686	609	609	613	616	616	lp_184	686	625	629	628	636	637
lp_208	702	633	619	622	635	635	lp_208	702	635	632	635	640	635
lp_251	2128	1856	1848	1864	1882	1883	lp_251	2128	1957	1941	1956	1950	1960
lp_324	1044	852	861	860	871	882	lp_324	1044	879	885	893	890	894

- Calculated upper bound using **Maximum Spanning Tree**
- Randomness can cause worse results even with more runtime
- Less improvements after 30 seconds



DETAILED RESULTS

Ip_172

My solution: 155
Exact: 164
Diff: 9

Ip_083

My solution: 164
Exact: 164

Ip_156

My solution: 260
Exact: 260

Ip_100

My solution: 194
Exact: 194

Ip_152

My solution: 230
Exact: 230

Ip_101

My solution: 205
Exact: 208
Diff: 3

Ip_136

My solution: 253
Exact: 253

Ip_175

My solution: 241
Exact: 241

Ip_135

My solution: 233
Exact: 236
Diff: 3

Ip_151

My solution: 264
Exact: 265
Diff: 1

Ip_108

My solution: 322
Exact: 322